

# CS657A: INFORMATION RETRIEVAL SCORED RETRIEVAL MODEL

Arnab Bhattacharya  
arnabb@cse.iitk.ac.in

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur  
<http://web.cse.iitk.ac.in/~cs657/>



2<sup>nd</sup> semester, 2021-22  
Tue 1030-1145, Thu 1200-1315

# Non-Boolean Match


- A query may be more general than just a set of terms
- It may itself be another document or some free text
- No document can then be expected to match it fully
- **Information retrieval task**
  - Given a free text query, find the most **similar** documents
- “Similarity” of a document  $d$  with the query  $q$  can be measured by a **score**  $s(d, q)$
- **Scoring information retrieval task**
  - Given a free text query, find the documents that are most similar, i.e., have the largest scores

# Zones

- Each document is generally associated with **metadata**, e.g., title, author, date, etc.
- Sometimes, queries on indexes of these fields, called **parametric indexes**, are also asked
  - Find documents where “Indian” appears in the title
- Fields are generalized to **zones** that may contain free text as well
- Separate inverted indexes can be built for each zone



- Or, zone may be mentioned *explicitly* in a single inverted index



# Weighted Zone Scoring

- Each zone has a weight that adds up to 1
  - For example, “title” has 0.3, and “body” has 0.7
- Given a query term, each zone scores a **zone score**
- It is 1 if the term is inside the zone; 0 otherwise
- **Weighted zone score** is the linear sum of the zone scores

$$s(q, d) = \sum_{\forall z \in d} s_z(q)$$

- Given a Boolean combination of query terms, each zone is scored
- These scores are then accumulated and ranked
- **Ranked Boolean retrieval**
- Weights of zones
  - Can be supplied by the application
  - Machine learned

# Term Frequency

- Moving away from the binary model
- If a document contains a query term more number of times, it is more important and should score higher
- Weight of a document  $d$  is, therefore, proportionate to the number of times the term  $t$  appears in it, called the **term frequency**

$$tf(t, d) = |t \in d|$$

- This assumes the **bag of words** model
- *Context* and *sequence* are lost
  - I love butter but I hate cheese
  - I love cheese but I hate butter

# Inverse Document Frequency

- Certain terms may appear across all or most documents
  - “bat” in cricket pages
- Consequently, they discriminate little among the documents and are not useful
- **Document frequency**,  $df_t$ , of a term  $t$  is the number of documents in the corpus that it appears in
  - Lesser is more discriminative
- Weight of a term is inversely proportionate to the documents it appears in
- If  $m$  is the total number of documents in the corpus

$$idf(t) = \log \frac{m}{df_t}$$

- This is called the **inverse document frequency**
- Logarithmic to make it less drastic
- Theoretical justification from the *log-odds* model
- Does not affect the ranking

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term  $t$  in a document  $d$  is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

- The **tf-idf** score has following properties
  - *Zero* if  $t$  does not appear at all in  $d$
  - *High* when  $t$  appears many times in  $d$
  - *Low* when  $t$  appears few times in  $d$
  - *High* when  $t$  appears in a small number of documents
  - *Low* when  $t$  appears in a large number of documents
- **Score** of a query  $q$  for a document  $d$  is

$$score(q, d) = \sum_{t \in q} tf-idf(t, d)$$

# Exercise

- $d_1$ : Water, water everywhere, not a drop to drink
- $d_2$ : I have filtered water
- $d_3$ : Don't drink and drive, rather drive and drink
- $d_4$ : Water quality is not good here
- $d_5$ : Milk is not good for health
- $d_6$ : Drinking water just after dinner is not healthy
- Query  $q$ : drinkable water
- Find tf, idf (with  $\log_2$ ) and tf-idf ( $\log_2$ ) scores, and rank

Terms	idf	tf						tf-idf					
		$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
drink	$\log_2 \frac{6}{3} = 1$	1	0	2	0	0	1	1	0	2	0	0	1
water	$\log_2 \frac{6}{4} = \frac{1}{2}$	2	1	0	1	0	1	1	$\frac{1}{2}$	0	$\frac{1}{2}$	0	$\frac{1}{2}$
Query	-	-	-	-	-	-	-	2	$\frac{1}{2}$	2	$\frac{1}{2}$	0	$\frac{3}{2}$

- Therefore, ranking is
  - $(d_1, d_3), (d_6), (d_2, d_4), (d_5)$



# Variants of Tf

- Tf is too drastic

$$tf(t, d) = \begin{cases} 1 + \log tf_{t,d} & \text{if present} \\ 0 & \text{otherwise} \end{cases}$$

- It still penalizes absent terms heavily

$$tf(t, d) = a + (1 - a) \frac{tf_{t,d}}{\max\{tf_{t,d}\}}$$

- $a$  is some default tf
- Very susceptible to outliers and stopwords
- Idf can be also modified as

$$idf(t) = \max \left\{ 0, \log \frac{m - df_t}{df_t} \right\}$$

# Document Vector

- Each document  $d$  has a score with each term  $t$  in the vocabulary
- Imagine an  $n$ -dimensional vector space where  $n$  is the total number of terms in the vocabulary
- Each document can be, thus, thought of as a vector (point) in this  $n$ -dimensional space
- Its coordinates are the scores corresponding to

$$d[t_i] = tf-idf(t_i, d)$$

- This is called the **document vector** model
- Document  $d$  is represented by its corresponding vector  $\vec{V}(d)$
- Longer documents have more number of terms and larger tf's
- To balance, document vectors can be *normalized* by their **length**

$$\vec{v}(d) = \frac{\vec{V}(d)}{|\vec{V}(d)|}$$

# Cosine Similarity

- Consider two documents  $d_1$  and  $d_2$  with their corresponding document vectors  $\vec{V}(d_1)$  and  $\vec{V}(d_2)$
- **Cosine similarity** measures the *normalized dot product*

$$\text{cosine-sim}(d_1, d_2) = \frac{\vec{V}(d_1) \cdot \vec{V}(d_2)}{|\vec{V}(d_1)| \cdot |\vec{V}(d_2)|}$$

- Measures the cosine of the angle between the vectors
- Consider the **length-normalized** document vectors
- Then, cosine similarity is their dot product

$$\text{co-sim}(d_1, d_2) = \vec{v}(d_1) \cdot \vec{v}(d_2)$$

## Example

Term	$d_1$	$d_2$	$d_3$
Indian	115	58	20
ancient	10	7	11
system	2	0	6
length	115.45	58.42	23.60

- Similarities between documents

$$\text{sim}(d_1, d_2) = \frac{115}{115.45} \cdot \frac{58}{58.42} + \frac{10}{115.45} \cdot \frac{7}{58.42} + \frac{2}{115.45} \cdot \frac{0}{58.42} = 0.99$$

$$\text{sim}(d_1, d_3) = \frac{115}{115.45} \cdot \frac{20}{23.60} + \frac{10}{115.45} \cdot \frac{11}{23.60} + \frac{2}{115.45} \cdot \frac{6}{23.60} = 0.88$$

$$\text{sim}(d_2, d_3) = \frac{58}{58.42} \cdot \frac{20}{23.60} + \frac{7}{58.42} \cdot \frac{11}{23.60} + \frac{0}{58.42} \cdot \frac{6}{23.60} = 0.89$$

- $d_1, d_2$  is the closest pair
- $d_3$  is closer to  $d_2$

# Query Vector



- Similar to a document, the query can also be viewed as a vector
- It is again just a *bag of words*
- Consider query  $q$  with the keywords “ancient” and “system”
- $\vec{V}(q) = (0, 1, 1)$  with  $\vec{v}(q) = (0.00, 0.71, 0.71)$
- Most similar document is the one with the *highest* cosine similarity

$$\arg \max_{d_i} \{\text{co-sim}(q, d_i)\} = \arg \max_{d_i} \{\vec{v}(q) \cdot \vec{v}(d_i)\}$$

- Document vector retrieval task
- If  $\vec{v}(q) \cdot \vec{v}(d_i)$  is the highest, then  $\vec{V}(q) \cdot \vec{V}(d_i)$  is the highest as well
  - Thus, normalization of query is not required

# Document Length Normalization

- Longer documents have more terms and, therefore, larger tf's
- Verbose documents may simply repeat terms to improve scores
- Documents on multiple topics need not be compensated
- Length normalization improves scores of shorter documents and deteriorates those of longer ones
- Ideally, with *known* queries and relevant documents



- Rotate about **pivot** to make cosine length similar

$$a \cdot |\vec{V}(d)| + (1 - a) \cdot l_p \approx a \cdot u_d + (1 - a) \cdot l_p$$

- $u_d$  is the number of unique terms in the document
- This is called **pivot length normalization**

# Improving Efficiency

- Brute-force method of computing scores with all the documents and ranking them is *not* scalable
- Only terms that appear in query need to be examined
  - Rest of the scores are 0
- Filter query terms whose *idf* is too low
  - Similar to stopwords
- *Pre-compute* **champion lists** for each term
  - Documents ranked for only that term
  - Offline process
  - Take union of top- $r$  of every query term to get top- $K$
- Build **tiered index**
  - Each level (tier) lists only those documents whose *tf* for the term is greater than a threshold
  - Continue with tiers till top- $K$  results are obtained

# Relevance Feedback

- Update query vector  $\vec{q}$  that maximizes similarity with relevant documents and minimizes that with irrelevant ones
- $D_r$  and  $D_n$  are sets of *actual* relevant and irrelevant documents

$$\vec{q}_{opt} = \arg \max_{\forall \vec{q}} [sim(\vec{q}, D_r) - sim(\vec{q}, D_n)]$$

- Using cosine similarity

$$\vec{q}_{opt} = \frac{1}{|D_r|} \sum_{\forall \vec{d}_r \in D_r} \vec{d}_r - \frac{1}{|D_n|} \sum_{\forall \vec{d}_n \in D_n} \vec{d}_n$$

- However,  $D_r$  and  $D_n$  are not known fully
- Only a partial subset  $U_r$  and  $U_n$  from user is known
- **Rocchio algorithm**

$$\vec{q}_m = \alpha \cdot \vec{q}_0 + \beta \cdot \frac{1}{|U_r|} \sum_{\forall \vec{d}_r \in U_r} \vec{d}_r - \gamma \cdot \frac{1}{|U_n|} \sum_{\forall \vec{d}_n \in U_n} \vec{d}_n$$

- Positive feedback is more important:  $\alpha = 1.00, \beta = 0.75, \gamma = 0.15$